

Breast Cancer Prediction Using Python

```
In [34]: ▶ # importing libraries
import numpy
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```


```
In [35]: ▶ # reading data from the file
df=pd.read_csv("data.csv")
```

```
In [36]: ▶ df.head()
```

Out[36]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothne
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

5 rows × 33 columns



In [37]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                 569 non-null    float64
12  radius_se                              569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                                569 non-null    float64
16  smoothness_se                          569 non-null    float64
17  compactness_se                         569 non-null    float64
18  concavity_se                           569 non-null    float64
19  concave points_se                      569 non-null    float64
20  symmetry_se                            569 non-null    float64
21  fractal_dimension_se                   569 non-null    float64
22  radius_worst                           569 non-null    float64
23  texture_worst                          569 non-null    float64
24  perimeter_worst                        569 non-null    float64
25  area_worst                             569 non-null    float64
26  smoothness_worst                       569 non-null    float64
27  compactness_worst                      569 non-null    float64
28  concavity_worst                        569 non-null    float64
29  concave points_worst                   569 non-null    float64
30  symmetry_worst                         569 non-null    float64
31  fractal_dimension_worst                 569 non-null    float64
32  Unnamed: 32                             0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

```
In [38]: # return all the columns with null values count  
df.isna().sum()
```

```
Out[38]: id                                0  
diagnosis                                0  
radius_mean                             0  
texture_mean                             0  
perimeter_mean                           0  
area_mean                                0  
smoothness_mean                          0  
compactness_mean                         0  
concavity_mean                           0  
concave points_mean                       0  
symmetry_mean                            0  
fractal_dimension_mean                    0  
radius_se                                 0  
texture_se                                0  
perimeter_se                              0  
area_se                                   0  
smoothness_se                             0  
compactness_se                            0  
concavity_se                              0  
concave points_se                         0  
symmetry_se                               0  
fractal_dimension_se                      0  
radius_worst                              0  
texture_worst                             0  
perimeter_worst                           0  
area_worst                                0  
smoothness_worst                          0  
compactness_worst                         0  
concavity_worst                           0  
concave points_worst                      0  
symmetry_worst                            0  
fractal_dimension_worst                   0  
Unnamed: 32                               569  
dtype: int64
```

```
In [39]: # return the size of dataset  
df.shape
```

```
Out[39]: (569, 33)
```

```
In [40]: # remove the column  
df=df.dropna(axis=1)
```

```
In [41]: # shape of dataset after removing the null column  
df.shape
```


```
Out[41]: (569, 32)
```

In [46]: `# describe the dataset`
`df.describe()`

Out[46]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.

8 rows × 7 columns




In [47]: `# Get the count of malignant<M> and Benign cells`
`df['diagnosis'].value_counts()`

Out[47]: B 357
M 212
Name: diagnosis, dtype: int64

In [48]: `sns.countplot(df['diagnosis'], label="count")`

/Applications/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[48]: <AxesSubplot:xlabel='diagnosis', ylabel='count'>



```
In [49]: # Label encoding(convert the value of M and B into 1 and 0)
from sklearn.preprocessing import LabelEncoder
labelencoder_Y = LabelEncoder()
df.iloc[:,1]=labelencoder_Y.fit_transform(df.iloc[:,1].values)
```

```
In [50]: df.head()
```

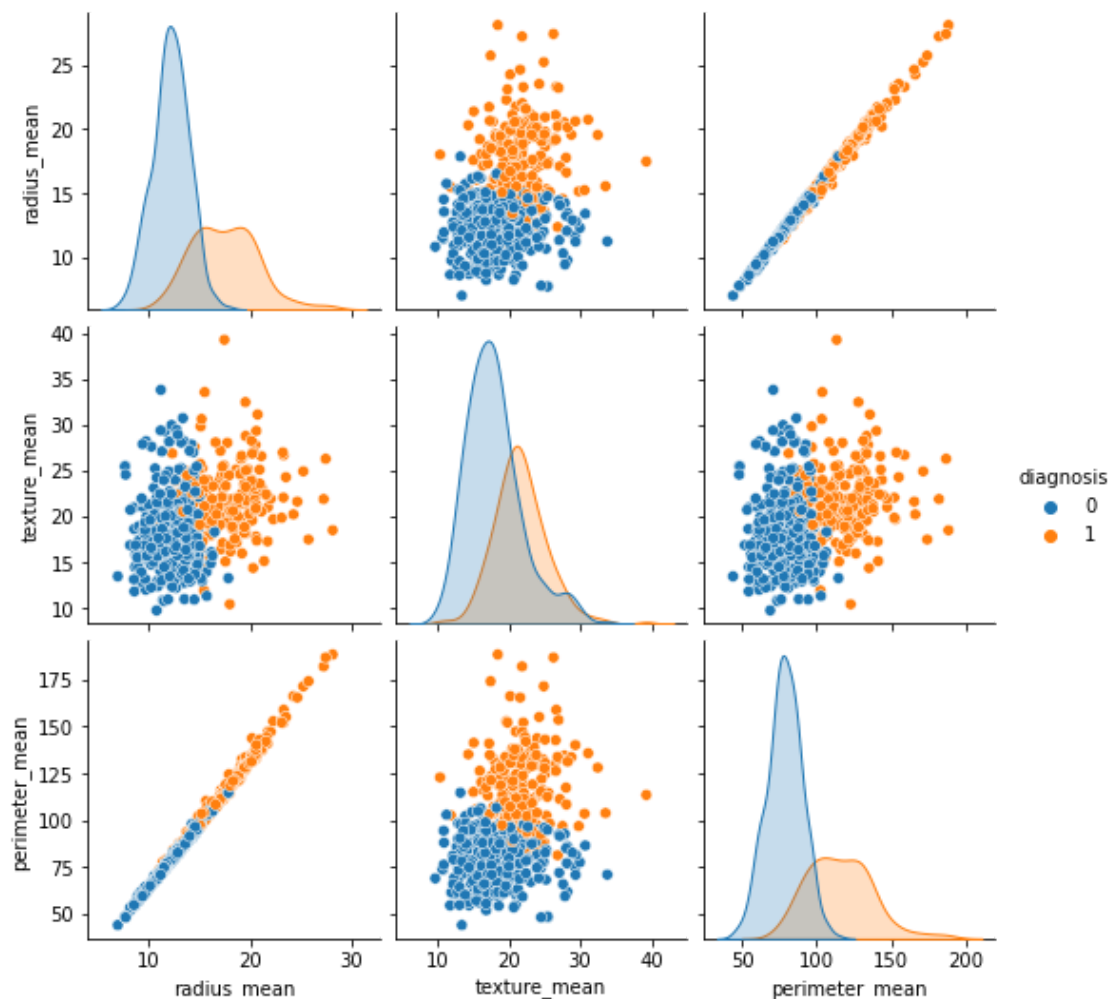
Out[50]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothne
0	842302	1	17.99	10.38	122.80	1001.0	
1	842517	1	20.57	17.77	132.90	1326.0	
2	84300903	1	19.69	21.25	130.00	1203.0	
3	84348301	1	11.42	20.38	77.58	386.1	
4	84358402	1	20.29	14.34	135.10	1297.0	

5 rows × 32 columns

```
In [16]: sns.pairplot(df.iloc[:,1:5],hue="diagnosis")
```

Out[16]: <seaborn.axisgrid.PairGrid at 0x7faeb2f65af0>




```
In [17]: # get the correlation
df.iloc[:,1:32].corr()
```

Out[17]:

ean	concave points_mean	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	area_wc
360	0.776614	0.330499	...	0.776454	0.456903	0.782914	0.7331
764	0.822529	0.147741	...	0.969539	0.297008	0.965137	0.9411
418	0.293464	0.071401	...	0.352573	0.912045	0.358040	0.3431
136	0.850977	0.183027	...	0.969476	0.303038	0.970387	0.9411
983	0.823269	0.151293	...	0.962746	0.287489	0.959120	0.9591
984	0.553695	0.557775	...	0.213120	0.036072	0.238853	0.2061
121	0.831135	0.602641	...	0.535315	0.248133	0.590210	0.5091
000	0.921391	0.500667	...	0.688236	0.299879	0.729565	0.6751
391	1.000000	0.462497	...	0.830318	0.292752	0.855923	0.8091
667	0.462497	1.000000	...	0.185728	0.090651	0.219169	0.1771
783	0.166917	0.479921	...	-0.253691	-0.051269	-0.205151	-0.2311
925	0.698050	0.303379	...	0.715065	0.194799	0.719684	0.7511
218	0.021480	0.128053	...	-0.111690	0.409003	-0.102242	-0.0831
391	0.710650	0.313893	...	0.697201	0.200371	0.721031	0.7301
427	0.690299	0.223970	...	0.757373	0.196497	0.761213	0.8111
564	0.027653	0.187321	...	-0.230691	-0.074743	-0.217304	-0.1821
279	0.490424	0.421659	...	0.204607	0.143003	0.260516	0.1991
270	0.439167	0.342627	...	0.186904	0.100241	0.226680	0.1881
260	0.615634	0.393298	...	0.358127	0.086741	0.394999	0.3421
009	0.095351	0.449137	...	-0.128121	-0.077473	-0.103753	-0.1101
301	0.257584	0.331786	...	-0.037488	-0.003195	-0.001000	-0.0221
236	0.830318	0.185728	...	1.000000	0.359921	0.993708	0.9841
879	0.292752	0.090651	...	0.359921	1.000000	0.365098	0.3451
565	0.855923	0.219169	...	0.993708	0.365098	1.000000	0.9771
987	0.809630	0.177193	...	0.984015	0.345842	0.977578	1.0001
822	0.452753	0.426675	...	0.216574	0.225429	0.236775	0.2091
968	0.667454	0.473200	...	0.475820	0.360832	0.529408	0.4381
103	0.752399	0.433721	...	0.573975	0.368366	0.618344	0.5431
323	0.910155	0.430297	...	0.787424	0.359755	0.816322	0.7471
464	0.375744	0.699826	...	0.243529	0.233027	0.269493	0.2091
930	0.368661	0.438413	...	0.093492	0.219122	0.138957	0.0791

```
In [51]: ▶ # visualize the correlation
plt.figure(figsize=(10,10))
sns.heatmap(df.iloc[:,1:10].corr(),annot=True,fmt=".0%")
```

Out[51]: <AxesSubplot:>



```
In [52]: ▶ # split the dataset into dependent(X) and Independent(Y) datasets
X=df.iloc[:,2:31].values
Y=df.iloc[:,1].values
```

```
In [53]: ▶ # splitting the data into training and test dataset
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.20,random_s
```

```
In [54]: ▶ # feature scaling
from sklearn.preprocessing import StandardScaler
X_train=StandardScaler().fit_transform(X_train)
X_test=StandardScaler().fit_transform(X_test)
```

```
In [55]: ▶ # models/ Algorithms

def models(X_train,Y_train):
    #Logistic regression
    from sklearn.linear_model import LogisticRegression
    log=LogisticRegression(random_state=0)
    log.fit(X_train,Y_train)

    #Decision Tree
    from sklearn.tree import DecisionTreeClassifier
    tree=DecisionTreeClassifier(random_state=0,criterion="entropy")
    tree.fit(X_train,Y_train)

    #Random Forest
    from sklearn.ensemble import RandomForestClassifier
    forest=RandomForestClassifier(random_state=0,criterion="entropy",n
    forest.fit(X_train,Y_train)

    print('[0]logistic regression accuracy:',log.score(X_train,Y_train)
    print('[1]Decision tree accuracy:',tree.score(X_train,Y_train))
    print('[2]Random forest accuracy:',forest.score(X_train,Y_train))

    return log,tree,forest
```

```
In [56]: ▶ model=models(X_train,Y_train)

[0]logistic regression accuracy: 0.9912087912087912
[1]Decision tree accuracy: 1.0
[2]Random forest accuracy: 0.9978021978021978
```


In [57]: ▶ *# testing the models/result*

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

for i in range(len(model)):
    print("Model",i)
    print(classification_report(Y_test,model[i].predict(X_test)))
    print('Accuracy : ',accuracy_score(Y_test,model[i].predict(X_test)))
```

Model 0

	precision	recall	f1-score	support
0	0.96	0.99	0.97	67
1	0.98	0.94	0.96	47
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Accuracy : 0.9649122807017544

Model 1

	precision	recall	f1-score	support
0	0.94	0.96	0.95	67
1	0.93	0.91	0.92	47
accuracy			0.94	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.94	0.94	0.94	114

Accuracy : 0.9385964912280702

Model 2

	precision	recall	f1-score	support
0	0.96	1.00	0.98	67
1	1.00	0.94	0.97	47
accuracy			0.97	114
macro avg	0.98	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

Accuracy : 0.9736842105263158

```
In [58]: ▶ # prediction of random-forest
pred=model[2].predict(X_test)
print('Predicted values:')
print(pred)
print('Actual values:')
print(Y_test)
```

Predicted values:

```
[1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 1 1 1 0 0 1 0 0 1 0 1 0 1 0 1
0
1 0 1 0 0 0 0 0 1 0 0 0 1 1 1 1 0 0 0 0 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0
0
1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 0 0 0 0 1 0 1 0 1 1
0
1 1 0]
```

Actual values:

```
[1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 1 1 1 0 0 1 0 0 1 0 1 0 1 0 1
0
1 0 1 1 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 0 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0
1
1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 0 0 0 0 1 0 1 0 1 1
0
1 1 0]
```

```
In [67]: ▶ plt.scatter()
plt.xlabel("actual ")
plt.ylabel("predicted")
plt.title("Actual VS Predicted")
plt.show()
```


```
-----
--
ValueError                                Traceback (most recent call last)
/var/folders/_1/z_zzbl013nb5zyqc3lvq4xh0000gn/T/ipykernel_1184/181843946
2.py in <module>
----> 1 plt.scatter(X,Y)
      2 plt.xlabel("actual ")
      3 plt.ylabel("predicted")
      4 plt.title("Actual VS Predicted")
      5 plt.show()

/Applications/anaconda3/lib/python3.9/site-packages/matplotlib/pyplot.py
in scatter(x, y, s, c, marker, cmap, norm, vmin, vmax, alpha, linewidths,
edgecolors, plotnonfinite, data, **kwargs)
   3066         vmin=None, vmax=None, alpha=None, linewidths=None, *,
   3067         edgecolors=None, plotnonfinite=False, data=None, **kwargs):
s):
-> 3068     __ret = gca().scatter(
   3069         x, y, s=s, c=c, marker=marker, cmap=cmap, norm=norm,
   3070         vmin=vmin, vmax=vmax, alpha=alpha, linewidths=linewidths,

/Applications/anaconda3/lib/python3.9/site-packages/matplotlib/__init__.p
y in inner(ax, data, *args, **kwargs)
   1359     def inner(ax, *args, data=None, **kwargs):
   1360         if data is None:
-> 1361             return func(ax, *map(sanitize_sequence, args), **kwargs)
   1362
   1363         bound = new_sig.bind(ax, *args, **kwargs)

/Applications/anaconda3/lib/python3.9/site-packages/matplotlib/axes/_axe
s.py in scatter(self, x, y, s, c, marker, cmap, norm, vmin, vmax, alpha,
linewidths, edgecolors, plotnonfinite, **kwargs)
   4496         y = np.ma.ravel(y)
   4497         if x.size != y.size:
-> 4498             raise ValueError("x and y must be the same size")
   4499
   4500         if s is None:
```

ValueError: x and y must be the same size



In []: ▶

In []: ▶